

Parallel Text Reference Breaker

Efficient Wikipedia Reference Detection using Aho-Corasick Algorithm and OpenMP

Rui Cai

Parallel Programming Course Project 2025

May 30, 2025

Outline

- 1 Project Overview
- 2 Usage and Applications
- 3 Algorithm and Data Structures
- 4 Parallel Implementation
- 5 Implementation Details
- 6 Results and Performance
- 7 Conclusion

Project Overview

- **Goal:** Automatically identify and annotate Wikipedia references in large text files
- **Approach:** Parallel text processing using Aho-Corasick string matching algorithm
- **Implementation:** C++11 with OpenMP parallelization
- **Output:** JSON format with reference annotations and positions

Key Features

- Offline Wikipedia titles database matching
- Parallel processing with configurable thread count
- Memory-efficient immutable trie structure
- 1KB block-based text segmentation


Problem Statement

Challenge:

- Large text files need semantic annotation
- Manual reference identification is time-consuming
- Need to match against millions of Wikipedia titles
- Performance requirements for real-time processing

Solution:

- Automated reference detection
- Parallel processing for scalability
- Efficient string matching algorithm
- Structured JSON output format



Usage Example

Command Line Interface

Basic usage

```
./text_reference_breaker wiki_titles.txt input.txt output.json
```

Sequential mode (for comparison)

```
./text_reference_breaker wiki_titles.txt input.txt output.json --sequential
```

Example with timing

```
time ./text_reference_breaker data/wiki_titles.txt large_text.txt result.json
```

JSON Output Format

```
{
  "text": "Coldplay had a tour concert in Helsinki in August, 2024.",
  "references": [
    { "start": 0, "end": 8, "titleIndex": 301, "title": "Coldplay" },
    { "start": 28, "end": 36, "titleIndex": 2105, "title": "Helsinki" },
    { "start": 40, "end": 46, "titleIndex": 984, "title": "August" }
  ]
}
```

Applications and Use Cases

Academic and Research

- Automatic citation generation for research papers
- Content analysis and topic modeling
- Knowledge graph construction

Content Management

- Wiki-style reference systems
- Educational content annotation
- News article fact-checking preparation

Data Processing

- Large-scale text corpus analysis
- Information extraction pipelines
- Semantic text enrichment

Aho-Corasick Algorithm

Time Complexity Analysis:

Construction Phase (Sequential)

- Build trie and failure links: $O(m)$
- Where $m = \sum_{i=0}^{|P|} |title_i|$ (total pattern length)
- Run once, immutable afterward


Search Phase (Parallelizable)

- Text traversal: $O(n + k)$
- Where n = text length, k = number of matches
- Each thread processes independent chunks

Patterns: "he", "she", "his", "hers"

Double circles = accepting states

Red dashed = failure links



Trie with failure links


Core Data Structures

```
struct Reference {  
    size_t start;           // Start position in text  
    size_t end;             // End position in text  
    size_t titleIndex;      // Index in Wikipedia titles  
    std::string title;      // Matched Wikipedia title  
};  
  
struct ProcessedBlock {  
    std::string text;       // Block content  
    std::vector<Reference> references; // Found references  
};  
  
class TextProcessor {  
private:  
    std::unique_ptr<aho_corasick::trie> trie_;  
    std::unordered_map<std::string, size_t> titleIndices_;
```


Parallelization Strategy

Block-Based Parallelization

- 1 **Text Segmentation:** Split input into 1KB blocks
- 2 **Parallel Processing:** Each thread processes independent blocks
- 3 **Result Aggregation:** Combine results maintaining order
- 4 **Position Adjustment:** Correct reference positions for global text



Listing 1: Parallel Processing Code

```
nlohmann::json TextProcessor::processFileParallel(const std::string& inputFile) {
    std::string text = utils::readFile(inputFile);
    std::vector<std::string> blocks = splitIntoBlocks(text);
    std::vector<ProcessedBlock> processedBlocks(blocks.size());

    size_t currentOffset = 0;

    #pragma omp parallel for schedule(dynamic)
    for (size_t i = 0; i < blocks.size(); ++i) {
        size_t blockOffset = 0;
        for (size_t j = 0; j < i; ++j) {
            blockOffset += blocks[j].length();
        }
        processedBlocks[i] = processBlock(blocks[i], blockOffset);
    }

    return blocksToJson(processedBlocks);
}
```

Key OpenMP Features:

- `#pragma omp parallel for` for loop parallelization
- `schedule(dynamic)` for load balancing
- Thread-safe trie operations (read-only after construction)

Performance Considerations

Memory Efficiency:

- Immutable trie structure
- Shared read-only data
- Minimal memory allocation per thread
- Block-based processing reduces memory footprint

Load Balancing:

- Dynamic scheduling
- Variable block processing times
- Automatic work distribution


Scalability Factors:

- Thread count = `omp_get_max_threads()`
- No thread synchronization during processing
- Linear speedup potential
- I/O bound for very large files

Cache Efficiency:

- Sequential block processing
- Locality of reference
- Trie fits in cache for common patterns

System Architecture



Key Components:

- **TextProcessor**: Main processing class
- **Aho-Corasick Trie**: Pattern matching engine
- **OpenMP**: Parallelization framework
- **JSON Output**: Structured result format

Build System and Dependencies

CMake Configuration

```
cmake_minimum_required(VERSION 3.10)
project(text_reference_breaker VERSION 1.0)
```

```
set(CMAKE_CXX_STANDARD 11)
find_package(OpenMP REQUIRED)
```

Dependencies:

- nlohmann/json (JSON processing)
- aho_corasick (Header-only string matching)
- OpenMP (Parallelization)

External Dependencies:

- **Aho-Corasick:** github.com/cjgdev/aho_corasick
- **nlohmann/json:** github.com/nlohmann/json
- **Test Data:** github.com/mmcky/nyu-econ-370

Performance Benchmarks

Test Environment

- **Hardware:** 16-core system with OpenMP
- **Compiler:** GCC with C++11 and OpenMP support
- **Wikipedia Databases:** 51 titles vs 10,000 titles
- **Input Texts:** 2.2KB sample vs 3.1MB War and Peace
- **Note:** Tables show processing time only (loading time: 6-13ms for 51 titles, 66-79ms for 10K titles)

Small Database (51 titles):

Input	Mode	Processing Time	Speedup
2.2KB	Sequential	5.64ms	-
2.2KB	Parallel	13.95ms	0.40x
3.1MB	Sequential	7863.76ms	-
3.1MB	Parallel	1296.74ms	6.06x

Large Database (10K titles):

Input	Mode	Processing Time	Speedup
2.2KB	Sequential	1252.71ms	-
2.2KB	Parallel	671.75ms	1.76x
3.1MB	Sequential	1766.43s	-
3.1MB	Parallel	285.80s	6.18x

Performance Analysis

Key Findings:

- **Small files:** Sequential faster (overhead dominates)
- **Large files:** Parallel provides 6x speedup consistently
- **Loading overhead:** 10K titles take 12x longer to load than 51 titles (66-79ms vs 6-13ms)
- **Processing complexity:** Large databases require 200x more processing time due to more matches
- **Extreme scaling:** 1766s \rightarrow 286s for largest test (6.18x speedup)

Scalability Analysis:

- **Consistent speedup:** 6x for large texts regardless of database size
- **Amdahl's Law:** Parallel portion dominates for large texts
- **Loading vs Processing:** Loading is one-time cost, processing scales with text size
- **Memory bandwidth:** Limits not reached even with 5M+ references
- **I/O impact:** Minimal compared to processing time

Why Not Use Full Wikipedia Title Dump?

Scale Challenge

- **Full Wikipedia dump:** 100MB+ title file
- **Current testing:** 51 titles (445 bytes) vs 10,000 titles (3.2MB)
- **Estimated full dump:** 300,000-350,000 titles (30-35x larger than 10K)
- **Time complexity:** Construction $O(m)$, Search $O(n + k)$

Why Not Use Full Wikipedia Title Dump?

Extrapolated Performance:

- **Loading time estimate:**
 - 10K titles: 70ms
 - 300K titles: 2-3 seconds
- **Processing time estimate:**
 - War & Peace (3.1MB text)
 - Sequential: 12-15 hours
 - Parallel: 2-2.5 hours

Practical Considerations:

- **Memory usage:** 100MB+ trie structure
- **Reference density:** 100M+ matches expected
- **Diminishing returns:** Over-annotation problem
- **Development focus:** Algorithm optimization vs data scale

Future Enhancements

Performance Improvements:

- GPU acceleration (CUDA/OpenCL)
- SIMD vectorization
- Memory-mapped I/O
- Distributed processing

Feature Extensions:

- **Fuzzy matching support**
- Multiple language support
- Real-time streaming processing
- **Web service API**

Algorithm Enhancements:

- Context-aware matching
- Machine learning integration
- Semantic similarity scoring
- Disambiguation algorithms

Scalability:

- Cloud deployment
- Microservice architecture
- Database integration
- Caching mechanisms

Acknowledgments

Open Source Libraries

- **Aho-Corasick Algorithm:** Header-only C++ implementation by cjddev
https://github.com/cjddev/aho_corasick
- **JSON Processing:** nlohmann/json library by Niels Lohmann
<https://github.com/nlohmann/json>
- **OpenMP:** Parallel programming API for shared memory multiprocessing

Data Sources

- **War and Peace Text:** Downloaded from NYU Economics 370 course materials
Maintained by mmcky at github.com/mmcky/nyu-econ-370
- **Google 10K English Words:** Common English vocabulary from
[first20hours/google-10000-english](https://first20hours.org/google-10000-english)
Based on n-gram frequency analysis of Google's Trillion Word Corpus
- **Wikipedia Titles:** Custom curated sample of 51 popular articles

Thank You!

Questions?

Project Repository: `github.com/RayChromium/parallel_text_breaker`

Technologies Used: C++11, OpenMP, Aho-Corasick, CMake

Key Metrics: Parallel processing, $O(n+m+z)$ complexity, 1KB block size